

Used Cars in Egypt

Import libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
```

```
In [ ]:
```

```
In [ ]:
```

Read and explore data

```
In [2]: cars_v1 = pd.read_csv('cars.csv')
```

```
In [4]: cars_v1.head()
```

```
Out[4]:
```

	Unnamed: 0	Brand	Model	Body	Color	Year	Fuel	Kilometers	Engine	Transmission	Price
0	5337	Hyundai	Accent	Sedan	Black	2007	Benzine	140000 to 159999	1600 CC	Automatic	140.0
1	5338	Hyundai	Accent	Sedan	Silver	2005	Benzine	180000 to 199999	1000 - 1300 CC	Manual	78.0
2	5339	Hyundai	Accent	Sedan	Gray	1999	Benzine	140000 to 159999	1400 - 1500 CC	Manual	70.0
3	5340	Hyundai	Accent	Sedan	Blue-Navy Blue	2009	Benzine	140000 to 159999	1600 CC	Automatic	150.0
4	5341	Hyundai	Accent	Sedan	Silver	2000	Benzine	10000 to 19999	1000 - 1300 CC	Manual	75.0

```
In [6]: cars_v1.describe(include='all')
```

Out[6]:

	Unnamed: 0	Brand	Model	Body	Color	Year	Fuel	Kilometers	Engine	Transn
count	14741.000000	14741	14741	14741	14741	14741.000000	14741	14741	14741	
unique	NaN	3	18	3	14	NaN	2	16	3	
top	NaN	Hyundai	128	Sedan	White	NaN	Benzine	More than 200000	1600 CC	
freq	NaN	5692	2425	13453	2614	NaN	14200	2505	6762	
mean	8934.846754	NaN	NaN	NaN	NaN	2005.456821	NaN	NaN	NaN	
std	4922.065495	NaN	NaN	NaN	NaN	12.655566	NaN	NaN	NaN	
min	812.000000	NaN	NaN	NaN	NaN	1970.000000	NaN	NaN	NaN	
25%	4497.000000	NaN	NaN	NaN	NaN	1998.000000	NaN	NaN	NaN	
50%	8182.000000	NaN	NaN	NaN	NaN	2010.000000	NaN	NaN	NaN	
75%	13373.000000	NaN	NaN	NaN	NaN	2015.000000	NaN	NaN	NaN	
max	17058.000000	NaN	NaN	NaN	NaN	2022.000000	NaN	NaN	NaN	

```
In [10]: for col in cars_v1.columns:
          print("Col is ", col)
          print(cars_v1[col].value_counts())
```

```

Col is Unnamed: 0
5337      1
2931      1
2933      1
2934      1
2935      1
..
16284     1
16285     1
16286     1
16287     1
14213     1

```

Name: Unnamed: 0, Length: 14741, dtype: int64

```

Col is Brand
Hyundai      5692
Fiat         5033
Chevrolet    4016
Name: Brand, dtype: int64

```

```

Col is Model
128         2425
Verna       1903
Elantra     1529
Lanos       1342
Accent      1272
Optra       1252
Shahin      1142
Aveo        994
131         572
Cruze       428
Uno         350
Avante      282
Tipo        274
Punto       270
Matrix      268
Tucson      182
I10         166
Excel       90
Name: Model, dtype: int64

```

```

Col is Body
Sedan       13453
Hatchback   1106
SUV         182
Name: Body, dtype: int64

```

```

Col is Color
White       2614
Black       2032
Silver      1952
Gray        1670
Red         1538
Blue- Navy Blue 1406
Other Color 1134
Burgundy    1061
Green       456
Gold        374
Beige       152
Brown       140
Yellow      134
Orange      78
Name: Color, dtype: int64

```

```

Col is Year

```

2013	850
2010	763
2011	728
2015	727
2012	693
2017	690
2014	622
2019	607
2009	591
2018	574
2016	562
2008	485
2021	388
2006	348
2020	346
2007	300
2001	260
1999	241
2000	231
2002	230
1998	219
2005	218
1987	212
2003	211
1990	211
1979	209
1982	198
2004	185
1997	183
1996	173
1980	172
1988	160
1983	160
1985	152
1981	145
1984	145
1993	144
1994	144
1986	143
1991	140
1977	130
2022	118
1989	115
1978	114
1995	114
1976	111
1992	90
1975	69
1974	56
1972	22
1973	21
1971	20
1970	1

Name: Year, dtype: int64

Col is Fuel

Benzine	14200
Natural Gas	541

Name: Fuel, dtype: int64

Col is Kilometers

More than 200000	2505
------------------	------

10000 to 19999	1666
180000 to 199999	1349
100000 to 119999	1192
0 to 9999	1088
140000 to 159999	1064
120000 to 139999	1005
90000 to 99999	996
160000 to 179999	760
20000 to 29999	612
80000 to 89999	560
50000 to 59999	436
60000 to 69999	402
40000 to 49999	372
30000 to 39999	370
70000 to 79999	364

Name: Kilometers, dtype: int64

Col is Engine

1600 CC	6762
1400 - 1500 CC	4356
1000 - 1300 CC	3623

Name: Engine, dtype: int64

Col is Transmission

Manual	9862
Automatic	4879

Name: Transmission, dtype: int64

Col is Price

115.0	254
23.0	234
138.0	209
161.4	195
185.6	195
...	
122.5	1
68.5	1
202.0	1
111.1	1
46.6	1

Name: Price, Length: 631, dtype: int64

Col is Gov

Cairo	4458
Giza	2412
Alexandria	1636
Sharqia	851
Qalyubia	806
Gharbia	630
Dakahlia	590
Monufia	444
Ismailia	360
Suez	308
Fayoum	276
Beheira	246
Minya	236
Asyut	216
Damietta	210
Beni Suef	186
Kafr al-Sheikh	174
Sohag	158
Red Sea	132
Port Said	128
Qena	100

```
South Sinai      56
Luxor            42
Aswan            42
Matruh           36
New Valley       8
Name: Gov, dtype: int64
```

```
In [ ]: cars_v1[['x', 'y']] = cars_v1['Kilometers'].str.extractall('(\d+)').unstack().loc[:,0]
```

```
In [19]: cars_v1['x']=cars_v1['x'].astype(int)
```

```
In [54]: cars_v1['y'].value_counts()
```

```
Out[54]: 19999      1666
         199999     1349
         119999     1192
         9999      1088
         159999     1064
         139999     1005
         99999     996
         179999     760
         29999     612
         89999     560
         59999     436
         69999     402
         49999     372
         39999     370
         79999     364
Name: y, dtype: int64
```

```
In [56]: cars_v1['y']=cars_v1['y'].astype(float)
```

```
In [ ]:
```

```
In [57]: cars_v1.dtypes
```

```
Out[57]: Unnamed: 0      int64
Brand      object
Model      object
Body       object
Color      object
Year       int64
Fuel       object
Kilometers object
Engine     object
Transmission object
Price      float64
Gov        object
x          int32
y          float64
KM_Adj     float64
dtype: object
```

```
In [58]: cars_v1['KM_Adj'] = cars_v1[['x', 'y']].mean(axis=1)
```

```
In [61]: cars_v1.drop(['Kilometers','x','y'],axis=1, inplace= True)
```

```
In [62]: cars_v1.head()
```

```
Out[62]:
```

	Unnamed: 0	Brand	Model	Body	Color	Year	Fuel	Engine	Transmission	Price	Gov	KM_A
0	5337	Hyundai	Accent	Sedan	Black	2007	Benzine	1600 CC	Automatic	140.0	Giza	149999
1	5338	Hyundai	Accent	Sedan	Silver	2005	Benzine	1000 - 1300 CC	Manual	78.0	Qena	189999
2	5339	Hyundai	Accent	Sedan	Gray	1999	Benzine	1400 - 1500 CC	Manual	70.0	Giza	149999
3	5340	Hyundai	Accent	Sedan	Blue-Navy Blue	2009	Benzine	1600 CC	Automatic	150.0	Cairo	149999
4	5341	Hyundai	Accent	Sedan	Silver	2000	Benzine	1000 - 1300 CC	Manual	75.0	Giza	149999



```
In [ ]:
```

```
In [63]: cars_v1[['x', 'y']] = cars_v1['Engine'].str.extractall('(\d+)').unstack().loc[:,0]
```

```
In [68]: cars_v1['x'] = cars_v1['x'].astype(int)
```

```
In [70]: cars_v1['y'] = cars_v1['y'].astype(float)
```

```
In [77]: cars_v1.loc[cars_v1['y'].isnull() , 'y'] = cars_v1['x']
```

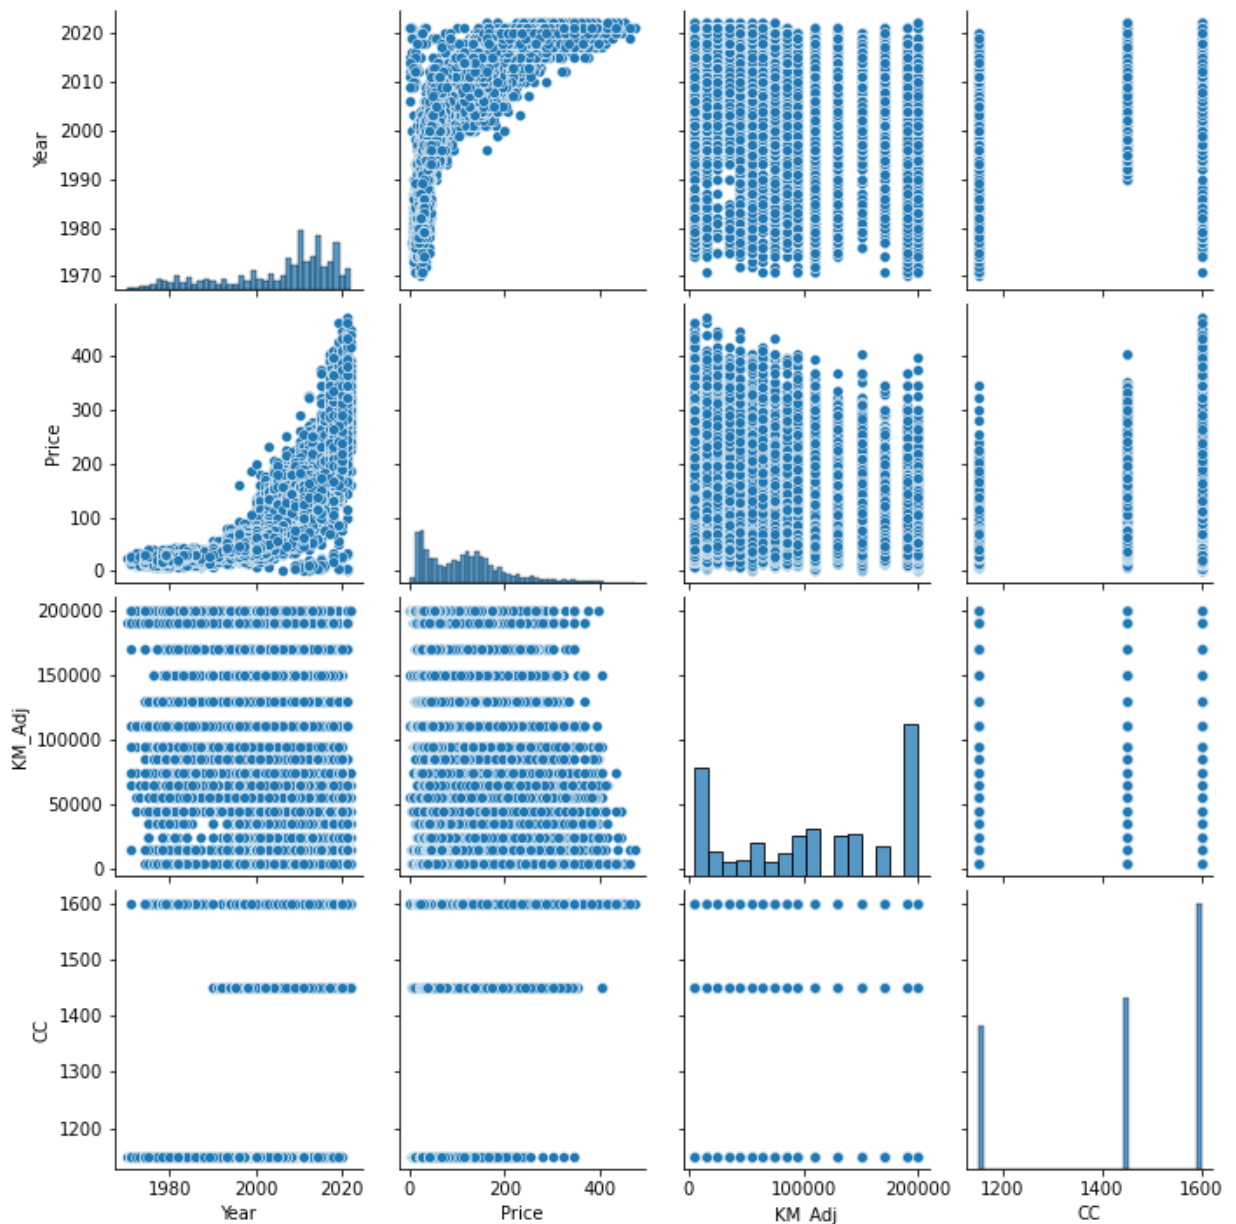
```
In [79]: cars_v1['CC'] = cars_v1[['x', 'y']].mean(axis=1)
```

```
In [81]: cars_v1.drop(['x', 'y'],axis =1 , inplace = True)
```

```
In [84]: cars_v1.drop(['Unnamed: 0'],axis =1 , inplace = True)
```


```
In [85]: sns.pairplot(cars_v1)
```

```
Out[85]: <seaborn.axisgrid.PairGrid at 0x18c03217fa0>
```



In [86]: `plt.hist(cars_v1['Price'])`


Out[86]: `(array([4098., 2394., 3274., 2762., 921., 604., 371., 204., 95., 18.]),
array([3. , 49.85, 96.7 , 143.55, 190.4 , 237.25, 284.1 , 330.95,
377.8 , 424.65, 471.5]),
<BarContainer object of 10 artists>)`



Starting the regression

```
In [88]: y = cars_v1['Price']
         x1 = cars_v1['Year']
```

```
In [89]: plt.scatter(y,x1)
         plt.xlabel('Price',fontsize= 20)
         plt.ylabel('Year',fontsize= 20)
         plt.show()
```



```
In [90]: x = sm.add_constant(x1)
         results = sm.OLS(y,x).fit()
         results.summary()
```

Out[90]:

OLS Regression Results


Dep. Variable:		Price		R-squared:		0.600
Model:		OLS		Adj. R-squared:		0.599
Method:		Least Squares		F-statistic:		2.206e+04
Date:		Thu, 10 Nov 2022		Prob (F-statistic):		0.00
Time:		12:48:49		Log-Likelihood:		-79165.
No. Observations:		14741		AIC:		1.583e+05
Df Residuals:		14739		BIC:		1.584e+05
Df Model:		1				
Covariance Type:		nonrobust				
	coef	std err	t	P> t	[0.025	0.975]
const	-9968.2243	67.894	-146.821	0.000	-1.01e+04	-9835.144
Year	5.0287	0.034	148.541	0.000	4.962	5.095
Omnibus:		3085.001	Durbin-Watson:		0.913	
Prob(Omnibus):		0.000	Jarque-Bera (JB):		7142.836	
Skew:		1.183	Prob(JB):		0.00	
Kurtosis:		5.455	Cond. No.		3.18e+05	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.18e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [91]: plt.scatter(x1,y)
y_hat = -9968.2243 + 5.0287* cars_v1['Year']
fig = plt.plot(x1, y_hat, lw=4, c= 'orange')
plt.ylabel('Price', fontsize= 20)
plt.xlabel('Year', fontsize= 20)
plt.show()
```



```
In [92]: y = cars_v1['Price']  
x1 = cars_v1[['Year', 'KM_Adj']]
```

```
In [93]: x = sm.add_constant(x1)  
results = sm.OLS(y,x).fit()  
results.summary()
```

Out[93]:

OLS Regression Results

Dep. Variable:	Price	R-squared:	0.601
Model:	OLS	Adj. R-squared:	0.601
Method:	Least Squares	F-statistic:	1.110e+04
Date:	Thu, 10 Nov 2022	Prob (F-statistic):	0.00
Time:	13:07:33	Log-Likelihood:	-79137.
No. Observations:	14741	AIC:	1.583e+05
Df Residuals:	14738	BIC:	1.583e+05
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-9842.2165	69.780	-141.045	0.000	-9978.995	-9705.438
Year	4.9684	0.035	143.121	0.000	4.900	5.036
KM_Adj	-4.757e-05	6.29e-06	-7.567	0.000	-5.99e-05	-3.52e-05

Omnibus:	2957.006	Durbin-Watson:	0.914
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6662.742
Skew:	1.148	Prob(JB):	0.00
Kurtosis:	5.361	Cond. No.	2.11e+07

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.11e+07. This might indicate that there are strong multicollinearity or other numerical problems.

Test OLS Assumptions

```
In [94]: cars_v1['cairo'] = cars_v1['Gov']
```

```
In [95]: cars_v1.loc[cars_v1['cairo'] == 'Cairo', 'cairo'] = 'Cairo'
```

```
In [96]: cars_v1.loc[cars_v1['cairo'] != 'Cairo', 'cairo'] = 'Not_Cairo'
```

```
In [97]: cars_v1['cairo'].value_counts()
```

```
Out[97]: Not_Cairo    10283
Cairo              4458
Name: cairo, dtype: int64
```

```
In [98]: cars_v1['cairo'] = cars_v1['cairo'].map({'Cairo':1, 'Not_Cairo':0})
```

```
In [99]: cars_v1['cairo'].value_counts()
```

```
Out[99]: 0    10283
         1     4458
         Name: cairo, dtype: int64
```

```
In [100]: y= cars_v1['Price']
          x1 = cars_v1[['Year', 'cairo']]
```

```
In [101]: x = sm.add_constant(x1)
          results = sm.OLS(y,x).fit()
          results.summary()
```

```
Out[101]:
```

OLS Regression Results						
Dep. Variable:	Price	R-squared:	0.600			
Model:	OLS	Adj. R-squared:	0.599			
Method:	Least Squares	F-statistic:	1.103e+04			
Date:	Thu, 10 Nov 2022	Prob (F-statistic):	0.00			
Time:	13:49:43	Log-Likelihood:	-79165.			
No. Observations:	14741	AIC:	1.583e+05			
Df Residuals:	14738	BIC:	1.584e+05			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
const	-9967.9346	67.902	-146.799	0.000	-1.01e+04	-9834.838
Year	5.0285	0.034	148.508	0.000	4.962	5.095
cairo	0.2941	0.933	0.315	0.753	-1.535	2.123
Omnibus:	3084.447	Durbin-Watson:	0.913			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7142.306			
Skew:	1.183	Prob(JB):	0.00			
Kurtosis:	5.455	Cond. No.	3.18e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.18e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Not Significant, Try another Factor

```
In [103]: cars_v1['CC'].value_counts()
```

```
Out[103]: 1600.0    6762
          1450.0    4356
          1150.0    3623
          Name: CC, dtype: int64
```

```
In [104]: cars_v1['new_CC'] = cars_v1['CC'].map({1600.0:1, 1450.0:0, 1150.0:0 })
```

```
In [106]: cars_v1['new_CC'].value_counts()
```

```
Out[106]: 0    7979
          1    6762
          Name: new_CC, dtype: int64
```

```
In [107]: y = cars_v1['Price']
          x1 = cars_v1[['Year', 'new_CC']]
```

```
In [108]: x = sm.add_constant(x1)
          results = sm.OLS(y,x).fit()
          results.summary()
```

```
Out[108]:
```

OLS Regression Results						
Dep. Variable:	Price	R-squared:	0.662			
Model:	OLS	Adj. R-squared:	0.662			
Method:	Least Squares	F-statistic:	1.441e+04			
Date:	Thu, 10 Nov 2022	Prob (F-statistic):	0.00			
Time:	13:53:26	Log-Likelihood:	-77923.			
No. Observations:	14741	AIC:	1.559e+05			
Df Residuals:	14738	BIC:	1.559e+05			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-8630.5232	67.498	-127.863	0.000	-8762.828	-8498.218
Year	4.3515	0.034	129.001	0.000	4.285	4.418
new_CC	44.5645	0.857	52.020	0.000	42.885	46.244
Omnibus:	3154.650	Durbin-Watson:	1.077			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	8964.914			
Skew:	1.128	Prob(JB):	0.00			
Kurtosis:	6.083	Cond. No.	3.44e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.44e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Car CC is significant, maybe without combining is better, but now we are trying the binary

```
In [109]: new_data = pd.DataFrame({'const':1, 'Year':[1990, 2015], 'new_CC':[0,1]})
new_data = new_data[['const', 'Year', 'new_CC']]
new_data
```

```
Out[109]:
```

	const	Year	new_CC
0	1	1990	0
1	1	2015	1


```
In [111]: predictions = results.predict(new_data)
predictions
```

```
Out[111]: 0      28.882579
1      182.233613
dtype: float64
```

```
In [124]: def pltcolor(lst):
cols=[]
for l in lst:
    if l==0:
        cols.append('red')
    elif l==1:
        cols.append('blue')
    return cols
# Create the colors list using the function above
cols=pltcolor(cars_v1['new_CC'])
```

```
In [128]: plt.yticks(np.arange(0, 500, step=20))
plt.scatter(cars_v1['Year'], cars_v1['Price'], alpha = .1, c=cols)
```

```
Out[128]: <matplotlib.collections.PathCollection at 0x18c0b764ac0>
```



Using SKLearn

```
In [129]: from sklearn.linear_model import LinearRegression
```

```
x = cars_v1['Year']  
y = cars_v1['Price']
```

x.shape, y.shape

 $((14741,), (14741,))$

```
x_matrix = x.values.reshape(-1,1)
```

```
reg = LinearRegression()
```

```
reg.fit(x_matrix,y)
```

LinearRegression()

```
reg.score(x_matrix, y)
```

0.5995221119507614

```
reg.coef_
```

```
array([5.02868432])
```

```
reg.intercept_
```

-9968.224279816577

```
reg.predict([[2023]])
```

```
array([204.80409602])
```

```
new_data = pd.DataFrame(data= [2000,2003,2005], columns =['Year'])
new_data
```

Year

0 2000

1 2003

2 2005

```
reg.predict(new_data)
```

```
D:\Data_Science\Anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but LinearRegression was fitted without feature names
```

```
warnings.warn(
```

```
array([ 89.1443567 , 104.23040966, 114.28777829])
```

```
new_data['Predected_Price']= reg.predict(new_data)
new_data
```



```
D:\Data_Science\Anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but LinearRegression was fitted without feature names
warnings.warn(
```

```
Out[156]:
```

	Year	Predicted_Price
0	2000	89.144357
1	2003	104.230410
2	2005	114.287778

```
In [157]: x = cars_v1[['Year', 'new_CC']]
          y = cars_v1['Price']
```

```
In [158]: reg = LinearRegression()
          reg.fit(x,y)
```

```
Out[158]: LinearRegression()
```

```
In [159]: reg.coef_
```

```
Out[159]: array([ 4.3514602 , 44.56452889])
```

```
In [160]: reg.intercept_
```

```
Out[160]: -8630.523223118082
```

Formula for adjusted R²

$$R_{adj.}^2 = 1 - (1 - R^2) * \frac{n-1}{n-p-1}$$

```
In [164]: r2 = reg.score(x,y)

          n = x.shape[0]

          p = x.shape[1]

          adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
          adjusted_r2
```

```
Out[164]: 0.6616009314286118
```

Feature selection

```
In [166]: from sklearn.feature_selection import f_regression
```

```
In [168]: f_regression(x,y)
          # 0 : f statistics
          # 1 : p value
          p_values = f_regression(x,y)[1]
          p_values.round(3)
          # P Value should be less than 5% to be significant
```

```
Out[168]: array([0., 0.])
```

Standrization

```
In [169]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()
```

```
In [170]: scaler.fit(x)
```

```
Out[170]: StandardScaler()
```

```
In [179]: x_scaled = scaler.transform(x)
```

Regression with Scalled features

```
In [180]: reg = LinearRegression()  
reg.fit(x_scaled, y)
```

```
Out[180]: LinearRegression()
```

```
In [182]: reg.coef_
```

```
Out[182]: array([55.06832355, 22.20619705])
```

```
In [183]: reg.intercept_
```

```
Out[183]: 116.5849874499691
```

```
In [186]: r2 = reg.score(x_scaled, y)  
n = x_scaled.shape[0]  
p = x_scaled.shape[1]  
  
adjusted_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)  
adjusted_r2
```

```
Out[186]: 0.6616009314286119
```

```
In [189]: reg_summary = pd.DataFrame([[ 'Intercept' ], [ 'Year' ], [ "new_CC" ]], columns=[ 'features' ])  
reg_summary[ 'weight' ] = reg.intercept_, reg.coef_[0], reg.coef_[1]
```

```
In [190]: reg_summary
```

```
Out[190]:
```

	features	weight
0	Intercept	116.584987
1	Year	55.068324
2	new_CC	22.206197

```
In [197]: new_data2 = pd.DataFrame({ 'Year': [1990, 2015, 1990, 2015], 'new_CC': [0, 1, 1, 0] })  
new_data2 = new_data2[[ 'Year', 'new_CC' ]]
```

```
new_data2
```

```
Out[197]:
```

	Year	new_CC
0	1990	0
1	2015	1
2	1990	1
3	2015	0

```
In [198]: reg.predict(new_data2)
```

D:\Data_Science\Anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but LinearRegression was fitted without feature names
warnings.warn(

```
Out[198]: array([109702.54885003, 111101.46313581, 109724.75504708, 111079.25693875])
```

```
In [199]: new_data2_scaled= scaler.transform(new_data2)
```

```
In [200]: reg.predict(new_data2_scaled)
```


```
Out[200]: array([ 28.88257921, 182.23361315,  73.44710809, 137.66908426])
```

```
In [ ]:
```

```
In [217]: sns.distplot(cars_v1['Price'])
```

D:\Data_Science\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[217]: <AxesSubplot:xlabel='Price', ylabel='Density'>
```




```
In [218]: price_log = np.log(cars_v1['Price'])
```

```
In [219]: sns.distplot(price_log)
```


```
D:\Data_Science\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[219]: <AxesSubplot:xlabel='Price', ylabel='Density'>
```




```
In [228]: plt.scatter(cars_v1['new_CC'], cars_v1['Price'])
```

```
Out[228]: <matplotlib.collections.PathCollection at 0x18c0e1abb20>
```




```
In [231]: plt.scatter(cars_v1['Year'], cars_v1['Price'])
```

```
Out[231]: <matplotlib.collections.PathCollection at 0x18c0e2d37c0>
```



```
In [230]: plt.scatter(cars_v1['Year'],price_log)
```

```
Out[230]: <matplotlib.collections.PathCollection at 0x18c0e18e130>
```



Multicollinearity

```
In [232]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [234]: cars_v1.columns.values
```

```
Out[234]: array(['Brand', 'Model', 'Body', 'Color', 'Year', 'Fuel', 'Engine',
        'Transmission', 'Price', 'Gov', 'KM_Adj', 'CC', 'cairo', 'new_CC'],
        dtype=object)
```

```
In [241]: variables = cars_v1[['Year', 'Price', 'KM_Adj', 'new_CC']]
```

```
In [242]: vif = pd.DataFrame()
```

```
In [243]: vif['VIF'] = [variance_inflation_factor(variables.values,i) for i in range(variables .sl
```

```
In [244]: vif['features'] = variables.columns
```

```
In [246]: #values between 1 and 5 are good
        # 10 is not acceptable
```

```
vif
```

```
Out[246]:
```

VIF features

0	6.782124	Year
1	4.369353	Price
2	3.582898	KM_Adj
3	2.564681	new_CC

```
In [251]:
```

```
cars_v1.describe(include='all')
```

```
Out[251]:
```

	Brand	Model	Body	Color	Year	Fuel	Engine	Transmission	Price	Gov
count	14741	14741	14741	14741	14741.000000	14741	14741	14741	14741.000000	14741
unique	3	18	3	14	NaN	2	3	2	NaN	2
top	Hyundai	128	Sedan	White	NaN	Benzine	1600 CC	Manual	NaN	Cairo
freq	5692	2425	13453	2614	NaN	14200	6762	9862	NaN	4451
mean	NaN	NaN	NaN	NaN	2005.456821	NaN	NaN	NaN	116.584987	NaN
std	NaN	NaN	NaN	NaN	12.655566	NaN	NaN	NaN	82.192718	NaN
min	NaN	NaN	NaN	NaN	1970.000000	NaN	NaN	NaN	3.000000	NaN
25%	NaN	NaN	NaN	NaN	1998.000000	NaN	NaN	NaN	43.700000	NaN
50%	NaN	NaN	NaN	NaN	2010.000000	NaN	NaN	NaN	110.000000	NaN
75%	NaN	NaN	NaN	NaN	2015.000000	NaN	NaN	NaN	161.000000	NaN
max	NaN	NaN	NaN	NaN	2022.000000	NaN	NaN	NaN	471.500000	NaN

Get Dummies

```
In [252]:
```

```
cars_dummies = pd.get_dummies(cars_v1,prefix=['Brand', 'Body'], columns=['Brand', 'Body
```

```
In [258]:
```

```
cars_dummies.drop(['cairo','new_CC'],axis= 1, inplace= True)
```

```
In [260]:
```

```
cars_dummies['log_price']= np.log(cars_dummies['Price'])
```

```
In [262]:
```

```
cars_dummies.columns
```

```
Out[262]:
```

```
Index(['Model', 'Color', 'Year', 'Fuel', 'Engine', 'Transmission', 'Price',  
      'Gov', 'KM_Adj', 'CC', 'Brand_Fiat', 'Brand_Hyundai', 'Body_SUV',  
      'Body_Sedan', 'log_price'],  
      dtype='object')
```

Linear Regression

```
In [264]:
```

```
targets = cars_dummies['log_price']
```

```
inputs = cars_dummies.drop(['Model', 'Color', 'Fuel', 'Engine', 'Transmission', 'Price',
```

```
In [265]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()
```

```
In [266]: scaler.fit(inputs)
```

```
Out[266]: StandardScaler()
```

```
In [267]: inputs_scaled = scaler.transform(inputs)
```

```
In [ ]:
```

Train Test Split

```
In [220]: from sklearn.model_selection import train_test_split
```

```
In [268]: X_train, X_test, y_train, y_test = train_test_split(inputs_scaled, targets, test_size=
```


```
In [269]: reg = LinearRegression()
```

```
In [270]: reg.fit(X_train, y_train)
```

```
Out[270]: LinearRegression()
```

```
In [272]: y_hat = reg.predict(X_train)
```

```
In [275]: plt.scatter(y_train, y_hat)  
plt.xlabel('Targets', fontsize= 20)  
plt.ylabel('Predictions', fontsize= 20)  
plt.show()
```




```
In [276]: sns.distplot(y_train - y_hat)
```

```
#Diff are normal with mean of zero
```

```
D:\Data_Science\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:
`distplot` is a deprecated function and will be removed in a future version. Please
adapt your code to use either `displot` (a figure-level function with similar flexibili
ty) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[276]: <AxesSubplot:xlabel='log_price', ylabel='Density'>
```



```
In [278]: reg.score(X_train, y_train)
```

```
Out[278]: 0.827772794127398
```

```
In [281]: reg.intercept_
```

```
Out[281]: 4.449310558911252
```

```
In [282]: reg.coef_
```

```
Out[282]: array([ 0.63839467, -0.02108024, -0.15173085,  0.04714382, -0.03700554,
                -0.05535697])
```

```
In [283]: reg_sum = pd.DataFrame(inputs.columns.values, columns = ['features'])
reg_sum['weights'] = reg.coef_
reg_sum
```

```
Out[283]:
```

	features	weights
--	----------	---------


0	Year	0.638395
1	KM_Adj	-0.021080
2	Brand_Fiat	-0.151731
3	Brand_Hyundai	0.047144
4	Body_SUV	-0.037006
5	Body_Sedan	-0.055357

Testing

```
In [285]: y_hat_test = reg.predict(X_test)
```



```
In [287]: plt.scatter(y_test,y_hat_test, alpha =.2)
plt.xlabel('Targets', fontsize= 20)
plt.ylabel('Predictions (y_hat_test)', fontsize= 20)
plt.show()
```



```
In [299]: df_pf = pd.DataFrame(np.exp(y_hat_test), columns =['Prediction'])
df_pf
```

```
Out[299]:
```

	Prediction
0	194.433879
1	179.646652
2	173.143321
3	95.367831
4	48.374679
...	...
2944	106.130016
2945	101.903494
2946	132.192594
2947	51.685648
2948	194.231250

2949 rows × 1 columns

```
In [300]: y_test= y_test.reset_index(drop=True)
```

```
In [301]: df_pf['Target'] = np.exp(y_test)
```

```
In [302]: df_pf.head()
```

Out[302]:

	Prediction	Target
0	194.433879	258.8
1	179.646652	212.8
2	173.143321	149.5
3	95.367831	130.0
4	48.374679	43.0

```
In [303]: df_pf['Residual'] = df_pf['Target'] - df_pf['Prediction']
```

```
In [304]: df_pf
```

Out[304]:

	Prediction	Target	Residual
0	194.433879	258.8	64.366121
1	179.646652	212.8	33.153348
2	173.143321	149.5	-23.643321
3	95.367831	130.0	34.632169
4	48.374679	43.0	-5.374679
...
2944	106.130016	135.0	28.869984
2945	101.903494	140.0	38.096506
2946	132.192594	85.0	-47.192594
2947	51.685648	46.0	-5.685648
2948	194.231250	264.5	70.268750

2949 rows × 3 columns

```
In [319]: df_pf['Diff%'] = np.absolute(df_pf['Residual']/df_pf['Target']*100).round(decimals=2)
df_pf
```

```
Out[319]:
```

	Prediction	Target	Residual	Diff%
0	194.433879	258.8	64.366121	24.87
1	179.646652	212.8	33.153348	15.58
2	173.143321	149.5	-23.643321	15.81
3	95.367831	130.0	34.632169	26.64
4	48.374679	43.0	-5.374679	12.50
...
2944	106.130016	135.0	28.869984	21.39
2945	101.903494	140.0	38.096506	27.21
2946	132.192594	85.0	-47.192594	55.52
2947	51.685648	46.0	-5.685648	12.36
2948	194.231250	264.5	70.268750	26.57

2949 rows × 4 columns

```
In [320]: df_pf.describe()
```

```
Out[320]:
```

	Prediction	Target	Residual	Diff%
count	2949.000000	2949.000000	2949.000000	2949.000000
mean	109.970049	119.555680	9.585631	31.897986
std	58.660083	83.287642	48.077432	147.508638
min	11.273800	3.000000	-220.185765	0.020000
25%	53.996347	46.000000	-16.945430	11.270000
50%	122.323310	115.000000	-0.258770	21.640000
75%	154.416791	161.400000	21.903266	34.860000
max	245.541000	460.000000	302.881023	4478.830000

```
In [ ]: #pd.options.display.max_rows = 999
        #pd.set_options('display.float_format', lambda x: "%.2f%" % x)
```

```
In [321]: f_regression(X_test,y_test)
```

```
Out[321]: (array([1.27346179e+04, 1.62972740e+02, 3.82886607e+03, 5.96788555e+02,
        1.75220808e+01, 3.98576639e+00]),
        array([0.00000000e+000, 2.28937348e-036, 0.00000000e+000, 3.50056158e-120,
        2.92302443e-005, 4.59780673e-002]))
```

```
In [ ]:
```

Using All Data

Get Dummies

```
In [329]: cars_v1.head()
cars_v2= cars_v1.drop(['CC', "cairo", 'new_CC'],axis=1)
```

```
In [330]:
```

```
In [331]: cars_v2['log_price']= np.log(cars_v2['Price'])
```

```
In [332]: cars_v2.drop(['Price'],axis= 1, inplace= True)
```

```
In [338]: cars_v2 = pd.get_dummies(cars_v2, drop_first=True)
```

```
In [ ]:
```

```
In [347]: targets = cars_v2['log_price']
inputs = cars_v2.drop(['log_price'], axis= 1)
```

```
In [348]: inputs.head()
```

```
Out[348]:
```

	Year	KM_Adj	Brand_Fiat	Brand_Hyundai	Model_131	Model_Accent	Model_Avante	Model_Aveo
0	2007	149999.5	0	1	0	1	0	0
1	2005	189999.5	0	1	0	1	0	0
2	1999	149999.5	0	1	0	1	0	0
3	2009	149999.5	0	1	0	1	0	0
4	2000	14999.5	0	1	0	1	0	0

5 rows × 65 columns


```
In [349]: scaler = StandardScaler()
scaler.fit(inputs)
inputs_scaled = scaler.transform(inputs)
```

```
In [350]: X_train, X_test, y_train, y_test = train_test_split(inputs_scaled, targets, test_size=
reg = LinearRegression()
reg.fit(X_train,y_train)
```

```
Out[350]: LinearRegression()
```

```
In [351]: y_hat = reg.predict(X_train)
```

```
In [352]: plt.scatter(y_train,y_hat)
plt.xlabel('Targets', fontsize= 20)
plt.ylabel('Predictions', fontsize= 20)
plt.show()
```



```
In [353]: reg.score(X_train, y_train)

reg.coef_

reg_sum = pd.DataFrame(inputs.columns.values, columns = ['features'])
reg_sum['weights'] = reg.coef_
reg_sum
```

```
Out[353]:
```


	features	weights
0	Year	4.969610e-01
1	KM_Adj	-1.682569e-03
2	Brand_Fiat	-4.429173e+10
3	Brand_Hyundai	1.400403e+10
4	Model_131	3.452594e-02
...
60	Gov_Red Sea	7.672787e-03
61	Gov_Sharqia	-1.606464e-03
62	Gov_Sohag	1.432896e-03
63	Gov_South Sinai	1.009703e-03
64	Gov_Suez	-2.652645e-03

65 rows × 2 columns

```
In [354]: y_hat_test = reg.predict(X_test)
```

```
In [355]: y_test= y_test.reset_index(drop=True)
```

```
In [356]: plt.scatter(y_test,y_hat_test, alpha =.2)
plt.xlabel('Targets', fontsize= 20)
plt.ylabel('Predictions (y_hat_test)', fontsize= 20)
plt.show()
```



```
In [357]: df_pf = pd.DataFrame(np.exp(y_hat_test), columns=['Prediction'])
df_pf
```

```
Out[357]:
```

	Prediction
--	------------

0	225.329036
1	211.546267
2	129.327479
3	114.999917
4	50.128709
...	...
2944	138.513329
2945	139.306055
2946	94.444913
2947	49.638902
2948	234.142508

2949 rows × 1 columns

```
In [358]: df_pf['Target'] = np.exp(y_test)
```

```
In [359]: df_pf.head()
```

```
Out[359]:
```

	Prediction	Target
0	225.329036	258.8
1	211.546267	212.8
2	129.327479	149.5
3	114.999917	130.0
4	50.128709	43.0

```
In [360]: df_pf['Residual']= df_pf['Target'] - df_pf['Prediction']

df_pf['Diff%'] = np.absolute(df_pf['Residual']/df_pf['Target']*100).round(decimals=2)
df_pf
```

```
Out[360]:
```

	Prediction	Target	Residual	Diff%
0	225.329036	258.8	33.470964	12.93
1	211.546267	212.8	1.253733	0.59
2	129.327479	149.5	20.172521	13.49
3	114.999917	130.0	15.000083	11.54
4	50.128709	43.0	-7.128709	16.58
...
2944	138.513329	135.0	-3.513329	2.60
2945	139.306055	140.0	0.693945	0.50
2946	94.444913	85.0	-9.444913	11.11
2947	49.638902	46.0	-3.638902	7.91
2948	234.142508	264.5	30.357492	11.48

2949 rows × 4 columns

```
In [373]: p_value_table = pd.DataFrame (data = inputs.columns, columns = ['Features' ])
```

```
In [375]: p_value_table['weights'] = f_regression(X_test,y_test)[1]
```

```
In [378]: p_value_table.sort_values('weights')
```

```
Out[378]:
```

	Features	weights
0	Year	0.000000e+00
2	Brand_Fiat	0.000000e+00
39	Transmission_Manual	3.854971e-315
38	Engine_1600 CC	2.967835e-198
3	Brand_Hyundai	3.500562e-120
...
24	Color_Blue- Navy Blue	8.544123e-01
62	Gov_Sohag	8.564710e-01
52	Gov_Luxor	8.804858e-01
49	Gov_Giza	9.314950e-01
48	Gov_Gharbia	9.540106e-01

65 rows × 2 columns

```
In [379]: # Gov is not Relevant, should be removed
```

```
In [381]: reg.score(X_test, y_test)
```

```
Out[381]: 0.9107993978077813
```

```
In [382]: # r2 for
```

```
r2 = reg.score(X_train,y_train)
n = X_train.shape[0]
p = X_train.shape[1]
adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
adjusted_r2
```

```
Out[382]: 0.9050039208760138
```

```
In [383]: r2 = reg.score(X_test,y_test)
```

```
n = X_test.shape[0]
p = X_test.shape[1]

adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
adjusted_r2
```

```
Out[383]: 0.9087882846816994
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```